

Credit EDA

September 22, 2017

Adriel Casellas

1 Credit Exploratory Data Analysis

1.0.1 Background

This dataset contains credit balance from several clients of a credit card company, along with several other characteristics of those clients. Credit card companies take losses when clients have too high, or too low of balances. Clients can declare bankruptcy when they are unable to pay their debts, but they allow for no accrued interest when the client's balance is too low. The most profitable clients are those with a moderate credit card balance upon which interest can be charged. Your task is to explore this data to find some of the relationships between the provided demographics and credit card balance.

1.0.2 Setup

```
In [107]: library(dplyr)
          library(ggplot2)
          library(cluster)
          library(ggdendro)

          credit <- read.csv('Credit.csv')
          glimpse(credit)
          summary(credit)
```

Observations: 310

Variables: 11

```
$ Income    <dbl> 14.891, 106.025, 104.593, 148.924, 55.882, 80.180, 20.996...
$ Limit     <int> 3606, 6645, 7075, 9504, 4897, 8047, 3388, 7114, 3300, 681...
$ Rating    <int> 283, 483, 514, 681, 357, 569, 259, 512, 266, 491, 589, 39...
$ Cards     <int> 2, 3, 4, 3, 2, 4, 2, 2, 5, 3, 4, 1, 1, 2, 3, 1, 2, 4, 1, ...
$ Age       <int> 34, 82, 71, 36, 68, 77, 37, 87, 66, 41, 30, 57, 49, 75, 6...
$ Education <int> 11, 15, 11, 11, 16, 10, 12, 9, 13, 19, 14, 7, 9, 13, 15, ...
$ Gender    <fctr> Male, Female, Male, Female, Male, Male, Female, Mal...
$ Student   <fctr> No, Yes, No, No, No, No, No, No, No, Yes, No, No, No, No...
$ Married   <fctr> Yes, Yes, No, No, Yes, No, No, No, No, Yes, Yes, Yes, Ye...
$ Ethnicity <fctr> Caucasian, Asian, Asian, Asian, Caucasian, Caucasian, Af...
```

```
$ Balance    <int> 333, 903, 580, 964, 331, 1151, 203, 872, 279, 1350, 1407,...
```

Income		Limit		Rating		Cards	
Min.	: 10.35	Min.	: 1160	Min.	:126.0	Min.	:1.000
1st Qu.:	23.15	1st Qu.:	3976	1st Qu.:	304.0	1st Qu.:	2.000
Median :	37.14	Median :	5147	Median :	380.0	Median :	3.000
Mean :	49.98	Mean :	5485	Mean :	405.1	Mean :	2.997
3rd Qu.:	63.74	3rd Qu.:	6453	3rd Qu.:	469.0	3rd Qu.:	4.000
Max.	:186.63	Max.	:13913	Max.	:982.0	Max.	:9.000

Age		Education		Gender	Student	Married
Min.	:23.00	Min.	: 5.00	Male :145	No :271	No :118
1st Qu.:	42.00	1st Qu.:	11.00	Female:165	Yes: 39	Yes:192
Median :	55.50	Median :	14.00			
Mean :	55.61	Mean :	13.43			
3rd Qu.:	69.00	3rd Qu.:	16.00			
Max.	:98.00	Max.	:20.00			

Ethnicity		Balance	
African American:	78	Min.	: 5.0
Asian	: 74	1st Qu.:	338.0
Caucasian	:158	Median :	637.5
		Mean :	671.0
		3rd Qu.:	960.8
		Max.	:1999.0

1.0.3 Analysis 1: Hclust

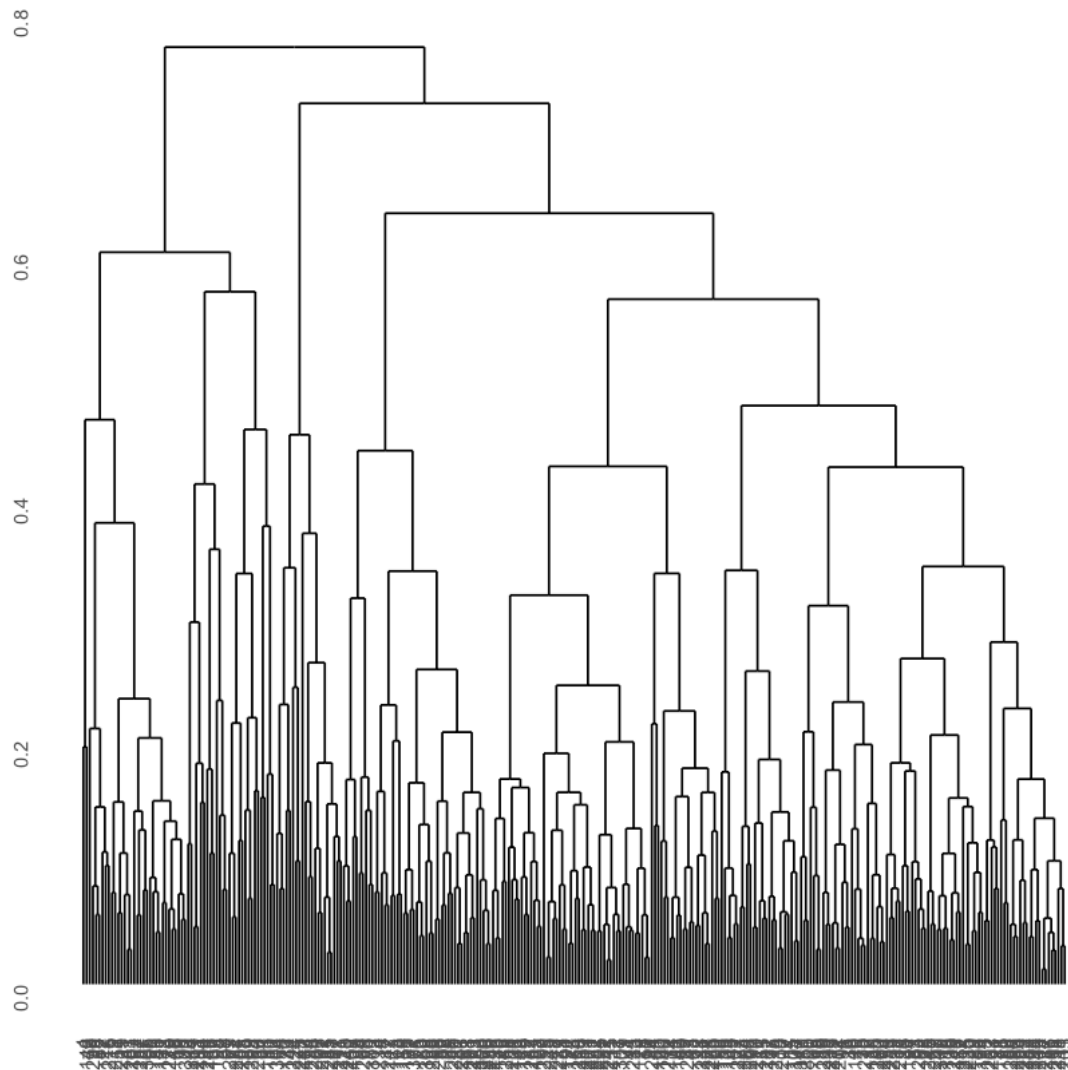
Create a distance matrix with the daisy function (this handles factor variables as well)

```
In [35]: credit.dist <- daisy(credit)
```

Invoke hclust() to perform hierarchical clustering on the distance matrix

```
In [36]: credit.hc <- hclust(credit.dist, method = "complete")
```

```
In [51]: gg dendrogram(credit.hc, k = 5, border="red")
```



Use cophenetic correlation coefficient to determine how well the dendrogram represents the distance matrix

```
In [52]: cor(cophenetic(credit.hc), credit.dist)
```

```
0.562325493005829
```

The CCPC indicates a mild fit to the distance matrix. Based on the groupings and the heights of the dendrogram groups, 4 clusters seems to sufficiently segment the data.

1.0.4 Analysis 2: Comparing differences between clusters

The cutree function returns the assignment vector for each observation

```
In [66]: credit.hc.seg <- cutree(credit.hc, k = 4)
```

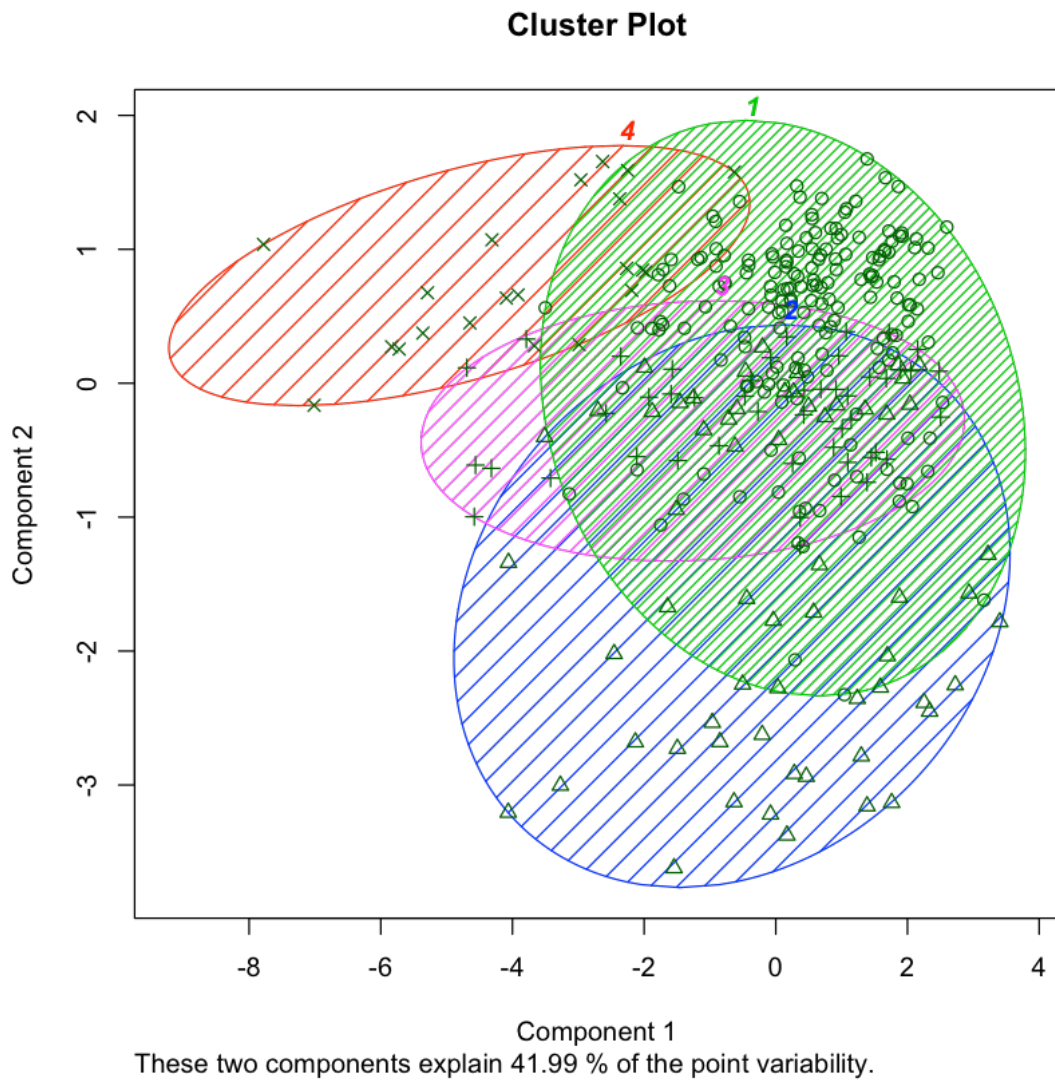
In order to compare differences between the segments, we will create a segment summary function to process the hclust output

```
In [67]: seg_summary <- function(data, cluster) {  
  # Ensure every variable is numeric  
  data_num <- data  
  for (i in 1:ncol(data)) {  
    data_num[,i] <- as.numeric(data[,i])  
  }  
  # Compute means by cluster/segment.  
  print(aggregate(data_num, list(cluster), mean))  
  clusplot(data_num, cluster, color=T, shade = T, labels=max(cluster), lines=0, main = )  
}
```

```
In [68]: seg_summary(credit, credit.hc.seg)
```

	Group.1	Income	Limit	Rating	Cards	Age	Education	Gender
1	1	40.45598	5006.227	372.7569	2.966851	55.36464	13.49724	1.607735
2	2	49.71639	5244.820	390.3770	3.000000	55.14754	13.13115	1.786885
3	3	53.25421	5753.383	419.7234	2.978723	55.38298	13.48936	1.000000
4	4	125.48814	9715.476	693.1905	3.285714	59.52381	13.52381	1.333333

	Student	Married	Ethnicity	Balance
1	1.016575	1.856354	2.165746	553.7238
2	1.590164	1.262295	2.540984	829.6230
3	1.000000	1.000000	2.170213	664.0213
4	1.000000	2.000000	2.428571	1236.4762



We will add the segment assignments to the original data frame to do some more exploratory analysis into the differences between the segments

```
In [93]: credit$Segment <- as.factor(credit.hc.seg)
```

```
In [106]: gg <- ggplot(credit, aes(x = Age, y = Income, col = Married))
           gg + geom_point() + geom_smooth(method = "lm") + facet_grid(. ~ Segment) + ggtitle(" ")
```



This facet grid shows distinct differences between the segments from the h-clustering. We will examine this in the summary

1.0.5 Summary of Analyses

The primary goal of this analysis was to look for any potential underlying groupings within the customer data that could help classify customers into certain groupings. The hierarchical clustering revealed that there are some underlying groupings that we could segment our customers by. After performing the H-clust, by looking at the heights of the clusters and the groupings, I concluded that 4 clusters would suffice for describing the underlying groupings. The CCPC of ~ 0.5 suggested that the dendrogram mildly represented the difference matrix of all the variables within the data set which is why I decided to use the dendrogram to choose and assign 4 assignment segments.

Looking closer at the clusters with `clusplot` revealed that the most distinguished groups (in terms of the 2 components that explain the most variability in the data) are groups 2 and 4. Overall, these groups had much in common, but all appeared to be distinguished enough to do some further analysis on to discover the differences between the groups. After observing the summary stats for all the groups (from the `seg_summary` function), these four groups can be classified as the following: 1) Lower-class families and singles, 2) Middle-class families and singles, 3) Single, social climbing men, 4) Upper-class married families. Identifying these groups allows for better service and better understanding of how customers needs can best be addressed.